

History of debuggers

Credits of slides belong to Levente Kurusa <levex@linux.com>
Imperial College London


Single user machines

- One of the first computers in the world
- Small application was loaded at the top of the memory
 - single step
 - examine registers
 - read/write memory

TX-0 at MIT



Batch processing machines



Debugged by putting macro call in the punch card and generating:

- Snapshots (*register dump*)
- Core dumps (*contents of memory*)

Then came CTSS (*Compatible Time-Sharing System*), one of the first time-sharing operating systems!

Debugging suddenly became interactive.

printf-debugging

```
*ptr = 1337;  
printf("Did we crash at line %d?\n", __LINE__);  
*((int *) 0) = 1337;  
printf("Did we crash at line %d?\n", __LINE__);
```

- The first version of UNIX had a debugger called, DB
- GNU had GDB and LLDB
- For Plan 9, ADB was created

These debuggers should be familiar!

Tracing processes

Most debuggers heavily rely on a system call known as ptrace.

The prototype of ptrace(2)

```
#include <sys/ptrace.h>
```

```
long ptrace(enum __ptrace_request request, pid_t pid,  
            void *addr, void *data);
```


How does `int a = 3, b = 0, c = a / b` result in a SIGFPE?

How does `int a = 3, b = 0, c = a / b` result in a SIGFPE?

1. Division by zero is noticed by the CPU

How does `int a = 3, b = 0, c = a / b` result in a SIGFPE?

1. Division by zero is noticed by the CPU
2. The CPU raises a divide-by-zero error (#DE)

How does `int a = 3, b = 0, c = a / b` result in a SIGFPE?

1. Division by zero is noticed by the CPU
2. The CPU raises a divide-by-zero error (#DE)
3. A handler in the kernel is eventually called

How does `int a = 3, b = 0, c = a / b` result in a SIGFPE?

1. Division by zero is noticed by the CPU
2. The CPU raises a divide-by-zero error (#DE)
3. A handler in the kernel is eventually called
4. The kernel sends a SIGFPE to the offending process

How does `int a = 3, b = 0, c = a / b` result in a SIGFPE?

1. Division by zero is noticed by the CPU
2. The CPU raises a divide-by-zero error (#DE)
3. A handler in the kernel is eventually called
4. The kernel sends a SIGFPE to the offending process
5. Your signal handler is called (or not if it is SIGKILL)

Implementation

- Enable tracing
- Run until system call
- Monitoring registers
- Single stepping
- Memory manipulation

Implementation

- Enable tracing
- Run until system call
- Monitoring registers
- Single stepping
- Memory manipulation

Implementation

- Enable tracing
 - `PTRACE_TRACEME`
- Run until system call
- Monitoring registers
- Single stepping
- Memory manipulation

Implementation

- Enable tracing
 - `PTRACE_TRACEME`
- Run until system call
- Monitoring registers
- Single stepping
- Memory manipulation

Implementation

- Enable tracing
 - `PTRACE_TRACEME`
- Run until system call
 - `PTRACE_SYSCALL`
 - `PTRACE_SYSEMU`
- Monitoring registers
- Single stepping
- Memory manipulation

Implementation

- Enable tracing
 - `PTRACE_TRACEME`
- Run until system call
 - `PTRACE_SYSCALL`
 - `PTRACE_SYSEMU`
- **Monitoring registers**
- Single stepping
- Memory manipulation

Implementation

- Enable tracing
 - `PTRACE_TRACEME`
- Run until system call
 - `PTRACE_SYSCALL`
 - `PTRACE_SYSEMU`
- **Monitoring registers**
 - `PTRACE_PEEKUSER` / `PTRACE_POKEUSER`
 - `PTRACE_GETREGS` / `PTRACE_SETREGS`
- Single stepping
- Memory manipulation

Implementation

- Enable tracing
 - `PTRACE_TRACEME`
- Run until system call
 - `PTRACE_SYSCALL`
 - `PTRACE_SYSEMU`
- Monitoring registers
 - `PTRACE_PEEKUSER` / `PTRACE_POKEUSER`
 - `PTRACE_GETREGS` / `PTRACE_SETREGS`
- Single stepping
- Memory manipulation

Implementation

- Enable tracing
 - `PTRACE_TRACEME`
- Run until system call
 - `PTRACE_SYSCALL`
 - `PTRACE_SYSEMU`
- Monitoring registers
 - `PTRACE_PEEKUSER` / `PTRACE_POKEUSER`
 - `PTRACE_GETREGS` / `PTRACE_SETREGS`
- Single stepping
 - `PTRACE_SINGLESTEP`
- Memory manipulation

Implementation

- Enable tracing
 - `PTRACE_TRACEME`
- Run until system call
 - `PTRACE_SYSCALL`
 - `PTRACE_SYSEMU`
- Monitoring registers
 - `PTRACE_PEEKUSER` / `PTRACE_POKEUSER`
 - `PTRACE_GETREGS` / `PTRACE_SETREGS`
- Single stepping
 - `PTRACE_SINGLESTEP`
- Memory manipulation

Implementation

- Enable tracing
 - `PTRACE_TRACEME`
- Run until system call
 - `PTRACE_SYSCALL`
 - `PTRACE_SYSEMU`
- Monitoring registers
 - `PTRACE_PEEKUSER` / `PTRACE_POKEUSER`
 - `PTRACE_GETREGS` / `PTRACE_SETREGS`
- Single stepping
 - `PTRACE_SINGLESTEP`
- **Memory manipulation**
 - `PTRACE_PEEKTEXT` / `PTRACE_POKETEXT`
 - `PTRACE_PEEKDATA` / `PTRACE_POKEDATA`

Architectural support

PTRACE_SINGLESTEP

PTRACE_SINGLESTEP

%EFLAGS

PTRACE_SINGLESTEP

%EFLAGS.TF

PTRACE_SINGLESTEP

Trap flag

(Sometimes referred to as "Trace flag")

PTRACE_SINGLESTEP

Trap flag *(Sometimes referred to as "Trace flag")*

- After each instruction, trap into #DB interrupt
- The kernel delivers a SIGTRAP
- This fact is delivered via a wait-event to the debugger

Breakpoints

Breakpoints

- `ud2` (machine code: `0x0F 0x0B`)
 - Triggers #UD – Undefined instruction exception

Breakpoints

- `ud2` (machine code: 0x0F 0x0B)
 - Triggers #UD – Undefined instruction exception
- `int $3` (machine code: 0xCC)
 - Triggers #BP – Breakpoint exception

Breakpoint implementation

- Replace with "int \$3"

Breakpoint implementation

- Replace with " `int` `$3`"
- Take note of the previous instruction

Breakpoint implementation

- Replace with "int \$3"
- Take note of the previous instruction
- When the breakpoint is hit, replace with the previous instruction

Breakpoint implementation

- Replace with "int \$3"
- Take note of the previous instruction
- When the breakpoint is hit, replace with the previous instruction
- Try executing the instruction again

Nota bene:

- Could have used "ud2"

- DR0 – DR3: Linear addresses

Debug registers

- DR0 – DR3: Linear addresses
- DR6: Debug control
 - Contains bitmasks for:
 - 1, 2, 4 or 8 bytes monitored
 - Break on read, write, execute, or read+write

Debug registers

- DR0 – DR3: Linear addresses
- DR6: Debug control
Contains bitmasks for:
 - 1, 2, 4 or 8 bytes monitored
 - Break on read, write, execute, or read+write
- DR7: Debug status
Bitmask showing which of DR0–DR3 triggered the #DB

Debug registers

- DR0 – DR3: Linear addresses
- DR6: Debug control
 - Contains bitmasks for:
 - 1, 2, 4 or 8 bytes monitored
 - Break on read, write, execute, or read+write
- DR7: Debug status
 - Bitmask showing which of DR0–DR3 triggered the #DB
- DR4 & DR5: Obsolete aliases to DR6 & DR7